

In [1]:
import pandas as pd
import warnings
warnings.filterwarnings("ignore")

In []:
student_df = pd.read_csv("/kaggle/input/student-performance-multiple-linear-regression/Student_Performance.csv")

In []:
student_df.head()

Out []:

	Hours Studied	Previous Scores	Extracurricular Activities	Sleep Hours	Sample Question Papers Practiced	Performance Index
0	7	99	Yes	9	1	91.0
1	4	82	No	4	2	65.0
2	8	51	Yes	7	2	45.0
3	5	52	Yes	5	2	36.0
4	7	75	No	8	5	66.0

In []:
student_df.shape

Out []:
(10000, 6)

In []:
student_df.info()

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10000 entries, 0 to 9999  
Data columns (total 6 columns):  
#   Column                                Non-Null Count  Dtype  
---  ---                                -  
0   Hours Studied                        10000 non-null  int64  
1   Previous Scores                     10000 non-null  int64  
2   Extracurricular Activities          10000 non-null  object  
3   Sleep Hours                         10000 non-null  int64  
4   Sample Question Papers Practiced    10000 non-null  int64  
5   Performance Index                   10000 non-null  float64  
dtypes: float64(1), int64(4), object(1)  
memory usage: 468.9+ KB
```

In []:
student_df.isnull().sum()

Out []:

```
Hours Studied      0  
Previous Scores    0  
Extracurricular Activities    0  
Sleep Hours        0  
Sample Question Papers Practiced    0  
Performance Index  0  
dtype: int64
```

In []:
student_df.duplicated().sum()

Out []:
127

In []:
student_df = student_df.drop_duplicates()

In []:
student_df.duplicated().sum()

Out []:
0

In []:
student_df.describe()

Out []:

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index
count	9873.000000	9873.000000	9873.000000	9873.000000	9873.000000
mean	4.992100	69.441102	6.531652	4.583004	55.216651
std	2.589081	17.325601	1.697683	2.867202	19.208570
min	1.000000	40.000000	4.000000	0.000000	10.000000
25%	3.000000	54.000000	5.000000	2.000000	40.000000
50%	5.000000	69.000000	7.000000	5.000000	55.000000
75%	7.000000	85.000000	8.000000	7.000000	70.000000
max	9.000000	99.000000	9.000000	9.000000	100.000000

In []:
clean_df = student_df.copy()
clean_df = pd.get_dummies(clean_df, columns = ['Extracurricular Activities'])

In []:
clean_df = clean_df.replace({True: 1, False: 0})
clean_df

Out []:

	Hours Studied	Previous Scores	Sleep Hours	Sample Question Papers Practiced	Performance Index	Extracurricular Activities_No	Extracurricular Activities_Yes
0	7	99	9	1	91.0	0	1
1	4	82	4	2	65.0	1	0
2	8	51	7	2	45.0	0	1
3	5	52	5	2	36.0	0	1
4	7	75	8	5	66.0	1	0
...
9995	1	49	4	2	23.0	0	1
9996	7	64	8	5	58.0	0	1
9997	6	83	8	5	74.0	0	1
9998	9	97	7	0	95.0	0	1
9999	7	74	8	1	64.0	1	0

9873 rows × 7 columns

In []:
X = clean_df.drop(['Performance Index'], axis = 1)
y = clean_df['Performance Index']

In []:
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 7)

In []:
from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
lin_reg.fit(X_train, y_train)

Out []:
LinearRegression

LinearRegression()

In []:
lin_pred = lin_reg.predict(X_test)
lin_pred

Out []:
array([60.70228604, 47.46488821, 47.61635294, ..., 61.69489914, 70.32626614, 44.30172536])

In []:
results_df = pd.DataFrame({'True Performance': y_test, 'Predicted Performance': lin_pred})
results_df

Out []:

	True Performance	Predicted Performance
3265	59.0	60.702286
175	45.0	47.464888
4225	48.0	47.616353
4653	64.0	65.771226
1321	24.0	20.788038
...
541	64.0	67.257694
5355	34.0	33.915870
2544	60.0	61.694899
3888	70.0	70.326266
1935	46.0	44.301725

1975 rows × 2 columns

In []:
lin_mae = mean_absolute_error(y_test, lin_pred)
lin_mae

Out []:
1.634425181582163

In []:
R2Score = r2_score(y_test, lin_pred)
print(f"R^2 Score is: {R2Score*100}")

R^2 Score is: 98.79370299201616

In []:
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

```
lin_pipeline = Pipeline([  
    ('scaler', StandardScaler()),  
    ('model', LinearRegression())  
)  
  
lin_pipeline.fit(X_train, y_train)  
lin_pipeline_pred = lin_pipeline.predict(X_test)
```

In []:
results_df = pd.DataFrame({'True Performance': y_test, 'Predicted Performance': lin_pipeline_pred})
results_df

Out []:

	True Performance	Predicted Performance
3265	59.0	60.715747
175	45.0	47.489653
4225	48.0	47.644918
4653	64.0	65.734780
1321	24.0	20.760702
...
541	64.0	67.281626
5355	34.0	33.932692
2544	60.0	61.722495
3888	70.0	70.313385
1935	46.0	44.318302

1975 rows × 2 columns

In []:
lin_pipeline_mae = mean_absolute_error(y_test, lin_pipeline_pred)
lin_pipeline_mae

Out []:
1.6354588740460978

In []:
R2Score = r2_score(y_test, lin_pipeline_pred)
print(f"R^2 Score is: {R2Score*100}")

R^2 Score is: 98.79291479120346